


One wire. A fresh newsroom each run.

Useful, timely Bitcoin news - with accountable sources and an independent edit. Autopost is OFF in this snapshot.



Continuity underneath

SQLite holds exact event identity, prior outputs, unresolved reporting work and NBN-native storylines. A fresh model context does not mean starting from zero.

Judgment, not quotas

The models decide importance, freshness, source sufficiency and framing. Code protects mechanical delivery integrity. Five to eight stories is an estimate, not a target to fill.

Many inputs. Different clocks.


The single worker sleeps 60 seconds after each cycle. Long network or model work can delay the next poll.

RSS + EDGAR	Each worker cycle	Reporting + upstream pilot feeds; Bitcoin-bearing 8-K filings. Haiku first.
X watches + guide accounts	180-second throttle	since_id recent search. Public list membership refreshes hourly.
Direct Perception	15-minute throttle	Broad Bitcoin discovery. Separate from the Node's provider usage.
Node wire API	5-minute consumer throttle	Supplemental versioned candidates. Valid pulse age: at most 3 hours.
Fresh EIC citations	Weekdays 14:40 / 21:15 UTC	Up to 12 cited reads per window; date and provenance checked.
Editorial desk	Normal 15-minute deadline	Nonempty eligible batches; priority, operator and backlog exceptions.

Node's last-known upstream pulse: hourly 05:00-20:00 Central. That scheduler lives elsewhere; NBN verifies payload age, not the scheduler itself. Scheduled Block threads are disabled.


Give the writer the right context.

No permanent conversation. One prepared batch, stable IDs, retrievable context and selective research.



From guide tip to reader-visible post.

Illustrative lifecycle, not a claim that any specific current story was published.



A future run can continue held work. A saved workbench is context, not a publishing command or a guaranteed retry clock.

What counts as proof? What is just context?

Source capability stays visible. A model's confident summary is not silently promoted to a captured source.

EVIDENCE

Inspect and attribute

Direct fetch: usable public source text, URL and fingerprint.

Native extract: source-specific, provider-reported paraphrase tied to an observed citation URL. Not verbatim page capture.

One credible report may support routine narrow news. Higher-stakes claims normally need stronger support or narrower attribution.

CONTEXT

Useful, but not proof

Guide posts, snippets, Node summaries, Luna prep and storylines direct attention.

A captured social post proves what its author said, not the underlying allegation.

The registry is guidance, not a closed universe. Syndicated copies do not become independent reports.

Exact event

Aliases + output lifecycle. Prevent a blind second create; preserve distinct developments.

Notebook + artifacts

30-day dated reporting memory. Partial work survives interruption; current publication wins.


Discoverable catalog

Notebooks, storylines and run artifacts. Search beyond Luna's suggestions; intake lookup covers 72h, up to 7d.

SHA-256 is a fingerprint for identity/integrity, not encryption. Model/provider credentials stay out of Desk snapshots and this guide.

Delivery is a lifecycle, not a button press.

Autopost OFF preserves the full editorial stack. It changes eligible delivery into a human Typefully draft.



The internal daily receipt audit is disabled by owner decision. Corrections remain human-reviewed. The rolling audit may turn autopost OFF for an evidenced systemic failure, never ON.

See the work. Know what the numbers mean.

A run-first workspace with a persistent wide-screen inspector. Visible browser tabs refresh every 15 seconds.

Newsroom	Latest run; Previous/Next pin history. Input, research, decisions, copy and activity.
Intake	Unique first-seen-day items, current status and preparation reasons. Search to drill down.
Research inspector	Assignments, returned findings, prepared captures and cited evidence. Missing is unknown.
Outputs	Locally tracked copy, receipt, timestamps and confirmation. Not every remote draft.
System & costs	Roster, stage/period costs, source health, PDF and paginated writer feedback.
Writer feedback	Optional self-report on each run and in System: what helped, what hindered, one suggestion. Unverified input for human discussion only; never injected into future writers or the editor.

History stays pinned until Back to latest. Rich observations expire after 14 days; gaps are not zero work. Editor fallbacks are not rewrites. Current publisher state has its own clock. Review tools retain guarded actions.

Observe, learn, make the smallest useful change.

The machinery now supports calibration; it does not replace editorial taste or the owner's authority.



COSTS

Measure the right bill

Recorded intake/prep/research/writer/editor usage includes tokens, cache shape, latency, native calls and cost provenance.

xAI reported totals where available; other supported calls are estimated.

Unknown costs are not zero.



AUDIT SCOPE

Observe and improve

Internal daily audit: disabled. Historical receipt checks remain available.

Codex rolling audit: health, quality, misses, peer speed, cost and smallest technical repairs.

Bounded approved-style prompt tuning is allowed; broader changes need approval.



RELEASE DISCIPLINE

Keep it reversible

Independent review when requested.

Offline tests, clean release archive, SQLite backup.

Deploy to the existing Railway worker; smoke and observe a natural cycle.

Restore the audit. Never enable autopost as a release step.

Cost boundaries

Direct per-run costs use distinct metered production runs; unlinked intake stays shared. Day/week/month averages use complete Central periods. Historical receipt audits, hosting, source/search services and Codex remain outside the ledger.

Current reference map: DOCUMENTATION.md. Detailed architecture: SYSTEM.md. Prompt authority: orientation-brief-v2.md. Historical plans and bake-offs remain dated evidence, not live instructions.